

Forecast Error & Loss Function

- Forecast Error :  $e = Y - \hat{Y}$
- Squared Loss :  $L(e) = (Y - \hat{Y})^2$
- Absolute Loss :  $L(e) = |Y - \hat{Y}|$
- Properties of loss functions :
- 1.  $L(0) = 0$
- 2.  $L(e) \geq 0, \forall e$
- 3. Non-increasing if  $e < 0$
- 4. Non-decreasing if  $e > 0$
- 5. Risk :  $R(Y) = E(L(e)) = E(L(Y - \hat{Y}))$

Bias Variance Trade Off

$E[(y - \hat{f}(x))^2] = Var(\hat{f}(x)) + [Bias(\hat{f}(x))]^2 + Var(\epsilon)$

Validation Set Approach

- 1. Split data into 2 sets (train and test)
- 2. Train the data and test it on the test set
- 3. Compute the MSE
- Cons :
- 1. loss of data  $\Rightarrow$  Higher MSE
- 2. estimate of the test error is highly variable

K-fold cross-validation

- 1. Split the data into K fold
- 2. For each degree, compute MSE for each k block
- 3. Average each MSE for each degree
- 1.  $CV(K) = \frac{1}{K} \sum_{k=1}^K MSE_k$
- 2.  $\hat{V}ar(MSE_k) = \frac{1}{K-1} \sum_{k=1}^K (MSE_k - \bar{MSE}_k)^2$
- 3.  $\hat{SE}(CV(K)) = \sqrt{\frac{1}{K} \hat{V}ar(MSE_k)}$
- 4.  $CV(K) = \frac{1}{K} \sum_{n_k} I(y_i \neq \hat{y}_i)$

Some Useful formulae

- 1.  $RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$
- 2.  $C_p = \frac{1}{n} (RSS + 2k\hat{\sigma}^2)$
- 3. Adjusted  $R^2 = 1 - \frac{n-1}{n-k-1} \frac{RSS}{TSS}$
- 4.  $AIC = \frac{1}{N\hat{\sigma}^2} (RSS + 2k\hat{\sigma}^2)$
- 5.  $BIC = \frac{1}{N\hat{\sigma}^2} (RSS + \log(N)k\hat{\sigma}^2)$

Best Subset Selection

- 1. Let  $\bar{p} \leq p$
- 2. For each  $k \in 0, \dots, \bar{p}$ , fit exactly k predictors to the model
- 3. Pick the highest  $R^2$  model for each k
- 4. Choose the best among these selected models using AIC, BIC, CV
- Cons :
- 1. High computational time,  $2^p$  possible models
- 2. The prediction is highly unstable
- 3. easy to overfit if P is large

Methods when applying IC

- 1. First approach : Use  $\hat{\sigma}^2$  directly from each model
- 2. Second approach : When  $P \ll N$ , use the  $\hat{\sigma}^2$  from the largest model
- 3. Third approach : When  $\frac{P}{N}$  is not small, use iterative procedure
- a. use  $Var(Y)$  as  $\sigma^2$  and select best model based on AIC/BIC
- b. from this selected model, take its  $\sigma^2$  and select best model based on AIC/BIC
- c. repeat this procedure until it converges

Forward Stepwise Selection

- 1. Starts from 1 regressor, fit all possible regressor to the model
- 2. Select the model with the highest Adjusted- $R^2$
- 3. Add regressors to the selected model, choose the one with the highest  $R^2$
- 4. Repeat the procedure

- 5. Choose the best model among the selected models for each degree using AIC/BIC/CV
- 6. Pros and Cons:
- a. P : Can be applied in high-dimensional scenarios, where  $n < p, \bar{p} < n$
- b. P : There are only  $1 + \frac{p(p+1)}{2}$  possible models for each p
- c. C : May not always find the best possible model

Backward Stepwise Selection

- 1. Exclude 1 regressor from the full model, choose the one with highest  $R^2$  (exclude the one with highest  $p - value$ )
- 2. Exclude 1 regressor from the selected model, choose the one with highest  $R^2$  (exclude the one with highest  $p - value$ )
- 3. Repeat the procedure
- 4. Choose the best among the selected model usign AIC/BIC/ $R^2$
- 5. Pros and Cons :
- a. P : Only p + 1 possible models for each p
- b. C : Constrain the search space to reduce variance, incur more bias
- c. C : Can't be use when  $P > N$

Ridge Regression

- 1. Incur some bias but greatly reduce variance
- 2.  $\sum_{i=1}^N (y_i - \beta_0 - \sum_{j=1}^P \beta_j x_{ij})^2 + \lambda \sum_{j=1}^P \beta_j^2$
- 3.  $||\beta||_2 = \sqrt{\sum_{j=1}^P \beta_j^2}$
- 4.  $\beta$  is not scale invariant, so we need to standardize variables before fitting them to the model
- 5.  $\hat{\beta}_\lambda^R = \frac{\hat{\beta}_{OLS}}{(1+\lambda)}$
- 6.  $E(\hat{\beta}_\lambda^R) = \frac{E(\hat{\beta}_{OLS})}{(1+\lambda)} = \frac{\beta}{(1+\lambda)}$
- 7.  $Var(\hat{\beta}_\lambda^R) = \frac{Var(\hat{\beta}_{OLS})}{(1+\lambda)}$
- 8. Estimators are biased but has lower variance than OLS
- 9. Pros and Cons :
- a. P : Can be used when  $P > N$
- b. P : Works best when some features are correlated
- c. C : Does not select the model simultaneously

How to choose  $\lambda$

- 1. Run the Ridge regression
- 2. Test different values for  $\lambda$ , use it to do CV
- 3. Choose the value of  $\lambda$  that minimizes MSE
- 4. higher  $\lambda$ , higher bias, lower variance

LASSO Regression

- 1. Shrink and select coefficients
- 2.  $\sum_{i=1}^N (y_i - \beta_0 - \sum_{j=1}^P \beta_j x_{ij})^2 + \lambda \sum_{j=1}^P |\beta_j|$
- 3.  $||\beta||_1 = \sum_{j=1}^P |\beta_j|$
- 4. Estimators are not scale invariant, need to standardize variables first
- 5. doesn't have closed form solution
- 6.  $\hat{\beta}_j^L(\lambda) = \hat{\beta}_j^{OLS} - \lambda/2$  if  $\hat{\beta}_j^{OLS} > \lambda/2$
- 7.  $\hat{\beta}_j^L(\lambda) = \hat{\beta}_j^{OLS} + \lambda/2$  if  $\hat{\beta}_j^{OLS} < -\lambda/2$
- 8.  $\hat{\beta}_j^L(\lambda) = 0$  if  $|\hat{\beta}_j^{OLS}| \leq -\lambda/2$
- 9. Pros and Cons :
- a. P : Helps to select variables
- b. P : Works even if  $P \gg N$
- c. C : Tends to select only one of the correlated features
- d. C : Can't use LOOCV because there is no closed form solution

How to choose  $\lambda$

- 1. Run CV many times and choose value of  $\lambda$  that occurs the most

- 2. 1SE rule : Calculate the SE for the estimated MSE for each  $\lambda$  on the grid and then select the largest  $\lambda$  for which the estimated CV MSE is within 1 SE of the minimum-MSE  $\lambda$  choice
- 3. Higher  $\lambda \Rightarrow$  lower penalization  $\Rightarrow$  lower model complexity
- 4. Plug-in  $\lambda$  :  $\lambda = 2c\sqrt{n}\hat{\sigma}\Phi^{-1}(1 - \gamma/2p)$ ,  $c = 1.1$  is suggested
- 5. Although LASSO selects variables, the coefficients are still biased
- 6. Sparsity (Some coefficient is 0 and non-zero coefficient is large) is assumed but may not be the case in reality

Post-LASSO Estimation

- 1. Run LASSO regression
- 2. Use the selected variables and run OLS again
- 3. Works well with theoretically motivated plug-in penalty ( $c = 1.1$  is suggested)
- 4. Sure Independence Screening (SIS) : rank variables by marginal correlations and only consider  $d$  most highly correlated to Y variables (i.e. run univariate regression on Y for each variable and choose  $d$  most correlated variables)
- 5.  $d = N - 1$  or  $d = N/\log(N)$  is recommended

Principal Components

- 1. Consider linear combinations  $Z_1, Z_2, \dots, Z_M$  of the original  $\mathbf{X}$
- 2.  $Z_m = \sum_{j=1}^P \phi_{jm} X_j$
- 3.  $y_i = \theta_0 + \sum_{m=1}^M \theta_m z_{im} + \epsilon_i$
- 4.  $\phi_1 = \text{argmax}_{\phi} \phi' = 1 \text{var}[\mathbf{X}\phi]$
- 5. where  $\phi_1$  is called loading/weight. What we are doing here is to find unit vectors that maximize  $\text{var}(\mathbf{X}\phi)$ , while they are orthogonal to each other.
- 6. PC,  $Z$ , is the linear combination of  $\mathbf{X}$ , where the coefficients are  $\phi$ , ranked by contribution to variance.
- 7.  $\phi$  are the eigenvectors of the covariance matrix  $\Sigma = \text{var}[\mathbf{X}]$
- 8.  $\text{var} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \Sigma = \begin{pmatrix} \sigma_1^2 & \sigma_{12} \\ \sigma_{21} & \sigma_2^2 \end{pmatrix}$  where  $\sigma_1^2 + \sigma_2^2 = \sigma_x^2$
- 9. An eigenvector  $u_j = (u_{j1} \quad u_{j2})'$  satisfied  $\Sigma u_j = D_j u_j$
- 10. Eigenvectors are orthogonal to each other (ie.  $u_1 \cdot u_2 = 0$ )
- 11.  $||u_1|| = \sqrt{u_{11}^2 + u_{12}^2}$
- 12.  $Z_i = \mathbf{X}\phi_1$ , where  $\phi_i$  is a unit eigenvector, and  $Z$  is a PC.
- 13.  $\text{var}(Z_i) = D_i$ , where  $D_i$  is  $\phi_i$ 's associated eigenvalue
- 14.  $Z$  is the linear combination of  $\mathbf{X}$  using eigenvectors as weights
- 15.  $\mathbf{X}$  can be expressed as a linear combination of  $Z$
- 16.  $\mathbf{X} = Z\Phi^{-1} = Z\Phi'$ , since  $\Phi$  is orthogonal matrix.
- 17. The ratio that gives the share of total variation in the data explained by each PC ( $Z_i$ ) is  $\frac{D_i}{\sum_{i=1}^T D_i}$
- 18.  $\Sigma = \begin{pmatrix} \phi_1 & \phi_2 \end{pmatrix} \times \begin{pmatrix} D_1 & 0 \\ 0 & D_2 \end{pmatrix} \times \begin{pmatrix} \phi_1' \\ \phi_2' \end{pmatrix} = D_1\phi_1\phi_1' + D_2\phi_2\phi_2'$
- 19. PC are not scale-invariant, we need to standardize the data
- 20. PCA works best on highly correlated data. If  $\mathbf{X}$  are uncorrelated, no dimension reduction via PCA is possible

PCR

- 1. Run PCA on the predictor matrix  $\mathbf{X}$
- 2. Construct the sample principal components  $Z$
- 3. Remove all but  $M < P$  first PC

- 4. Regress  $y$  on  $z_1, \dots, z_M$
- 1. It can be shown that Ridge shrinks the coefficients of PC by  $\frac{D_j}{D_j + \lambda}$
- 2. Thus, PC corresponding to lowest variance directions get shrunk the most
- 3. Assumption: the directions of highest variability in  $\mathbf{X}$  are those most associated with  $y$ , but this may not be the case
- 4. There is no variable selection in PCR, all  $\mathbf{X}$  are used to form PC

PLS

- 1. PLS finds directions helpful in explaining both  $\mathbf{Y}$  and  $\mathbf{X}$
- 2. To find the direction  $Z_1$ , regress  $\mathbf{Y}$  on each individual  $\mathbf{X}_j$  and set the weights  $\phi_{1j}$  to the respective OLS coefficients
- 3. We estimate  $y_i = \phi_{1p}x_{ip} + \epsilon_i, \forall p$
- 4. Then construct  $Z_1 = \sum_{j=1}^P \phi_{1j}X_j$ , this will be our first PC
- 5. After that,  $x_{ip} = \beta_p Z_{1i} + u_{pi}, \forall p$
- 6.  $y_i = \phi_{2p}\hat{u}_{ip} + \epsilon_i, \forall p$
- 7. Then the second PC is  $Z_2 = \sum_{j=1}^P \hat{\phi}_{2j}\hat{u}_j$
- 8. Repeat the steps above

Regression Trees

- 1. Regression trees splits the covariate space into a set of rectangles and then fit a simple model in each one.
- 2.  $f(x) = \sum_{m=1}^M c_m I(x \in R_m)$
- 3.  $\hat{c}_m = \text{ave}(y_i | x_i \in R_m)$
- 4. Ideally, we want to minimize  $\sum_{m=1}^M \sum_{i \in R_m} (y_i - \hat{y}_{R_m})^2$ , but this is not possible
- 5. STEP 1: define  $R_1(j, s) = X | X_j < s$  and  $R_2(j, s) = X | X_j \geq s$  and choose  $j$  and  $s$  as:  $\text{argmin} \sum_{i: x_i \in R_1(j, s)} (y_i - \hat{y}_{i, R_1})^2 + \sum_{i: x_i \in R_2(j, s)} (y_i - \hat{y}_{i, R_2})^2$
- 6. We do argmin for each dimension, and for each dimension, do argmin to find  $s$ .
- 7. The set of leaves gives a partition of the X space
- 8. Continuous data, the prediction will be the sample mean of Y in that node
- 9. Categorical data, the prediction will be the majority in that node
- 10. Trees handle both categorical and numeric X and Y well. We no need to worry about the scale of the X. It is computationally efficient. Small trees are very interpretable. Performs automatic interaction detection. Does variable selection.
- 11. Does not handle categorical variables when there are a lot of categories. Deep trees lose interpretability. Unstable with respect to training data.
- 12. Tree choice can be represented as a minimization problem:  $T_\alpha = \arg \min_T Q(T, y) = \arg \min_T L(T, y) + \alpha |T|$
- 13. In quadratic loss,  $T_\alpha = \arg \min_T \sum_{m=1}^M \sum_{i: x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T|$
- 14.  $\alpha$  is the penalty parameter. If  $\alpha = 0$ , we would choose the tree that perfectly fits the data resulting in overfitting.

Classification Trees

- 1. Loss function:  $L(T, y) = \sum_{m=1}^T N_m L_m(T, y_m)$ ,  $m$  indexes terminal nodes in tree T.  $N_m$  is the no of obs. in terminal node  $m$ .  $y_m$  is the vector of outcomes for obs. in terminal node  $m$ .
- 2.  $\hat{p}_{mk} = \frac{1}{N_m} \sum_{i: x_i \in R_m} 1(y_i = k)$ ,  $k(m) = \arg \max_k \hat{p}_{mk}$
- 3. Misclassification error:

$\frac{1}{N_m} \sum_{i \in R_m} 1(y_i \neq k(m)) = 1 - \hat{p}_{mk}(m)$  4. Gini index:  
 $\sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk})$   
completely pure  $\Rightarrow$  Gini index = 0  
5. Cross-entropy/deviance:  
 $-\sum_{k=1}^K \hat{p}_{mk} \log(\hat{p}_{mk})$   
6. We use gini index or deviance to grow the tree, and use misclassification error to prune the tree.

Performance Measures

- 1. Lift =  $\frac{TP}{FP+FP}$ , within positive, how many are correct.
- 2. ROC =  $\frac{TP}{FP}$ , we want to maximize AUC.
- 3. Both lift and ROC consider performance of  $\hat{p}(x)$  by looking at predictions  $\hat{y} = 1(\hat{p}(x) > s)$  as  $s$  varies
- 4.  $s = 1 \Rightarrow 1(\hat{p}(x)) = 0 \Rightarrow \hat{y}$  predicts all negative
- 4.  $s = 0 \Rightarrow 1(\hat{p}(x)) = 1 \Rightarrow \hat{y}$  predicts all positive

Bagging

- 1. We fit a lot of high variance low bias models from bootstrap samples. Then we average them together to reduce variance.
- 2.  $\hat{f}_B \text{ag}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^b(x)$
- 3. Typically choose B large enough for the procedure to settle down.
- 4. We don't need to do CV. We will automatically have B/3 out of bag observations.

Random Forests

- 1. Bagging trees are correlated because data are drawn from the same sample.
- 2. RF does not consider all P available predictors for each tree. This decorrelates trees.
- 3. In each bootstrap draw, use only m j P predictors.  $m = \sqrt{P}$  for classification and  $m = P/3$  for regression problems.
- 4. Bagging is RF when m = P

Boosting

- 1. Fit lots of small, low variance high bias models, aggregate them slowly to improve forecasting
- 2. Basis function expansions take the form  $f(x) = \sum_{m=1}^M \beta_m b(x; \gamma_m)$ ,  $b(x; \gamma)$  are simple functions,  $\gamma$  are parameters in that simple function
- 3. These models are fit by minimizing a loss function averaged over the training data  
 $\min_{\beta, \gamma} \sum_{i=1}^N L(y_i, f_{m-1} + \beta b(x_i; \gamma))$

FSAM

- 1. Initialize  $f_0(x) = 0$
- 2. For m = 1 to M:
  - a. compute:  $\text{argmin}_{\beta, \gamma} \sum_{i=1}^N L(y_i, f_{m-1} + \beta b(x_i; \gamma))$
  - b. set  $f_m(x) = f_{m-1}(x) + \beta_m b(x, \gamma_m)$
  - 3. The final model is just  $f(x) = f_M(x) = \sum_{m=1}^M \beta_m b(x, \gamma_m)$

Gradient Boosting

- 1. Initialize  $f_0(x) = \text{argmin}_{\gamma} \sum L(y_i, \gamma)$  for a suitable loss function
- 2. For m = 1 to M:
  - a. compute  $r_{im} = -\frac{\partial L(y, f(x))}{\partial f(x)}$
  - b. fit a base learner  $h_m(x)$  to  $r_{im}$
  - c. compute coefficient for  $h_m(x)$  as:  
 $\gamma_m = \text{argmin}_{\gamma} \sum_{i=1}^N L(y_i, f_{m-1}(x_i) + \gamma h_m(x_i))$
  - d. update model to  $f_m(x) = f_{m-1}(x) + \lambda \gamma_m h_m(x)$  for  $\lambda \leq 1$
  - 3. The final model is just  $f_M(x)$

Boosting Regression Trees

- 1. Initialize  $f_0(x) = 0$
- 2. For m = 1 to M:
  - a. compute  $r_{im} = y_i - f_{m-1}(x_i)$
  - b. fit a regression tree to residuals
  - c. for each node, compute  
 $\gamma_{jm} = \text{argmin}_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma)$   
 $\gamma_{jm} = \frac{\sum r_{im}}{\text{Number of observations in } R_{jm}}$
  - d. update:  
 $f_m(x) = f_{m-1}(x) + \lambda \sum_{j=1}^{J_m} \gamma_{jm} 1(x \in R_{jm})$
  - 3. The prediction model is then  $f_M(x)$

Loss for K-class Classification

- 1. Multinomial logit:  
 $Pr(Y = k|X = x) = \frac{\exp(f_k(x))}{\sum_{j=1}^K \exp(f_j(x))}$
- 2. Multinomial deviance:  
 $L(y, p(x)) = -\sum_{k=1}^K 1(y = k) f_k(x) + \log(\sum_{j=1}^K \exp(f_j(x)))$

Gradient Boosting for K-class problems

- 1. Initialize  $f_{k0}(x) = 0$
- 2. For m = 1 to M:
  - a. set  $p_k(x) = \frac{\exp(f_{k, m-1}(x))}{\sum_{j=1}^K \exp(f_{j, m-1}(x))}$

- b. for k = 1,..., K
  - i. compute  $r_{ikm} = y_{ik} - p_k(x_i)$
  - ii. fit a tree to  $r_{ikm}$   
 $\gamma_{jkm} = \frac{K-1}{K} \frac{\sum_{x_i \in R_{jkm}} r_{ikm}}{\sum_{x_i \in R_{jkm}} |r_{ikm}|(1-|r_{ikm}|)}$
  - iii.  $\gamma_{jkm} = \frac{K-1}{K} \frac{\sum_{x_i \in R_{jkm}} r_{ikm}}{\sum_{x_i \in R_{jkm}} |r_{ikm}|(1-|r_{ikm}|)}$
  - iv. update  
 $f_{km}(x) = f_{k, m-1}(x) + \lambda \sum_{j=1}^{J_m} \gamma_{jkm} 1(x \in R_{jkm})$
- 3. Output  $\hat{f}_k(x) = f_{km}(x)$

Relative Importance

- 1.  $J_l^2(T) = \sum_{t=1}^{J-1} \frac{1}{t} 1(v(t) = l)$   
This is the importance of  $X_l$  in a tree.
- 2.  $J_l^2 = \sum_{m=1}^M J_l^2(T_m)$   
This is the importance of  $X_l$  in a forest.
- 3. Average importance across trees in a forest:  
 $J_{\mathbb{K}}^2 = \frac{1}{M} \sum_{m=1}^M J_{lk}^2(T_{km})$
- 4. Average importance across forests:  
 $J_l^2 = \frac{1}{K} \sum_{k=1}^K J_{\mathbb{K}}^2$

ANN

- 1. The central idea is to extract linear combinations of the inputs as derived features, and then model the outcome as a nonlinear target of these features.
- 2. NN is a fancy nonlinear regression model which can be neatly represented by network diagrams.  
 $Z_m = \sigma(\alpha_{0m} + \alpha_m^T X)$ , linear comb. of X into activation function  
 $T_k = \beta_{0k} + \beta_k^T Z$ , linear comb. of Z  
 $f_k(X) = g_k(T)$ , output function from T

Activation Functions

- 1.  $\sigma(v) = 1/(1 + e^{-v})$
- 2.  $\sigma(z) = \max\{0, z\}$

Output Functions

- 1. For regression, usually is the identity function:  
 $g_k(T) = T_k$
- 2. For classification, usually the softmax function:  
 $g_k(T) = \frac{e^{T_k}}{\sum_{l=1}^K e^{T_l}}$

Fitting NN

- 1. For regression, the measure of fit is the squared loss:  
 $R(\theta) = \sum_{i=1}^N (y_{ik} - f_k(x_i))^2$
- 2. For classification, the measure of fit is deviance:  
 $R(\theta) = -\sum_{i=1}^N \sum_{k=1}^K y_{ik} \log f_k(x_i)$

Let  $z_{mi} = \sigma(\alpha_{0m} + \alpha_m^T x_i)$   
 $R(\theta) = \sum_{i=1}^N \sum_{k=1}^K (y_{ik} - f_k(x_i))^2$   
 $\frac{\partial R_i}{\partial \beta_{km}} = -2(y_{ik} - f_k(x_i))g'_k(\beta_k^T z_i)z_{mi} = \delta_{ki}z_{mi}$   
 $\frac{\partial R_i}{\partial \alpha_{ml}} = -\sum_{k=1}^K 2(y_{ik} - f_k(x_i))g'_k(\beta_k^T z_i)\beta_{km}\sigma'(\alpha_m^T x_i)x_{il} = s_{mi}x_{il}$   
 $\beta_{km}^{(r)}r + 1) = \beta_{km}^{(r)} - \gamma_r \sum_{i=1}^N \frac{\partial R_i}{\partial \beta_{km}^{(r)}}$   
 $\alpha_{ml}^{(r)}r + 1) = \alpha_{ml}^{(r)} - \gamma_r \sum_{i=1}^N \frac{\partial R_i}{\partial \alpha_{ml}^{(r)}}$

- $s_{mi} = \sigma'(\alpha_m^T x_i) \sum_{k=1}^K \beta_{km} \delta_{ki}$
- 1. Forward pass: the current weights are fixed and the predictions  $\hat{f}_k(x_i)$  are computed from the network.
- 2. Backward pass: the errors  $\delta_{ki}$  are computed, and then back-propagated via the equation above to get the errors  $s_{mi}$ .
- 3. Then, both sets of errors are used to compute the gradients for the updating steps.

Dealing with Overfit

- 1. Early stopping: stop before the global minimum.
- 2. Regularization:  $R(\theta) + \lambda J(\theta)$ , where  
 $J(\theta) = \sum_{k,m} \beta_{km}^2 + \sum_{m,l} \alpha_{ml}^2$

Input Scaling

- 1. ANN is sensitive to the scaling of the inputs
- 2. It is best to normalize all predictors

Hybrid Learning

- 1. Firstly, fit LASSO and produce  $\hat{f}(x)_{lasso}$
- 2. Secondly, compute the residuals  
 $r_i = Y_i - \hat{f}(x_i)_{lasso}$  and train the random forest on this residuals using the same predictors.
- 3. Form the hybrid prediction:  
 $\hat{f}(x) = \hat{f}(x)_{lasso} + \hat{f}(x)_{rf}$

Ensemble Learning

$f(x) = \sum_{k=1}^K w_k f_k(x)$   
 $f = w f_1 + (1 - w) f_2$   
 $Var(f) = w^2 \sigma_1^2 + (1 - w)^2 \sigma_2^2$   
 $w^* = \frac{\sigma_2^2}{\sigma_1^2 + \sigma_2^2} = \frac{\sigma_1^{-2}}{\sigma_1^{-2} + \sigma_2^{-2}}$   
 $Y_i = \beta_1 f_{1i} + \beta_2 f_{2i} + \dots + \beta_K f_{Ki} + e_i$   
 $Min(RSS + \lambda \sum_{i=1}^K |\beta_i|)$